

CO2 - AT2

1. Employee Data: Preprocessing, Aggregation & Visualization

PYTHON

Code

```
import matplotlib
matplotlib.use("Agg")
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns

# 1. GENERATE RAW MOCK DATA (Simulating an imperfect CSV file)
raw_data = {
    "Employee_Name": [" Alice ", "Bob", "Charlie", "Diana", "Evan"],
    "Department": ["HR", "Engineering", "HR", "Engineering", "Marketing"],
    "Salary": [50000, 85000, np.nan, 92000, 60000], # Missing value
    "Join_Date": [
        "2022-01-15",
        "2021-06-20",
        "2023-03-11",
        "2020-11-01",
        "2024-02-28",
    ],
}

df = pd.DataFrame(raw_data)
print("--- Raw Data ---")
print(df, "\n")

# 2. DATA PREPROCESSING
# Clean whitespace from text strings
df["Employee_Name"] = df["Employee_Name"].str.strip()

# Handle missing values by replacing NaN with the median salary
median_salary = df["Salary"].median()
df["Salary"] = df["Salary"].fillna(median_salary)

# Convert string dates into proper datetime format
df["Join_Date"] = pd.to_datetime(df["Join_Date"])

# Create a calculated feature (Years of service based on current year 2026)
df["Years_of_Service"] = 2026 - df["Join_Date"].dt.year

print("--- Preprocessed Data ---")
print(df, "\n")

# 3. DATA AGGREGATION
# Group by department and compute average metrics
dept_summary = (
    df.groupby("Department")
    .agg(Avg_Salary=("Salary", "mean"), Total_Employees=("Employee_Name", "count"))
    .reset_index()
)

print("--- Aggregated Summary ---")
print(dept_summary, "\n")

# 4. DATA VISUALIZATION
```

```

sns.set_theme(style="whitegrid")
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

sns.barplot(
    data=dept_summary,
    x="Department",
    y="Avg_Salary",
    hue="Department",
    ax=axes[0],
    palette="Blues_d",
    legend=False,
)
axes[0].set_title("Average Salary by Department")
axes[0].set_xlabel("Department")
axes[0].set_ylabel("Salary ($)")

sns.scatterplot( data=df,
    x="Years_of_Service",
    y="Salary",
    hue="Department",
    style="Department",
    s=200,
    ax=axes[1],
)
axes[1].set_title("Salary vs. Years of Service")
axes[1].set_xlabel("Years of Service")
axes[1].set_ylabel("Salary ($)")

plt.tight_layout()
plt.savefig("s1_plot.png", dpi=130)
plt.close()

```

Output

--- Raw Data ---


	Employee_Name	Department	Salary	Join_Date
0	Alice	HR	50000.0	2022-01-15
1	Bob	Engineering	85000.0	2021-06-20
2	Charlie	HR	NaN	2023-03-11
3	Diana	Engineering	92000.0	2020-11-01
4	Evan	Marketing	60000.0	2024-02-28

--- Preprocessed Data ---

	Employee_Name	Department	Salary	Join_Date	Years_of_Service
0	Alice	HR	50000.0	2022-01-15	4
1	Bob	Engineering	85000.0	2021-06-20	5
2	Charlie	HR	72500.0	2023-03-11	3
3	Diana	Engineering	92000.0	2020-11-01	6
4	Evan	Marketing	60000.0	2024-02-28	2

--- Aggregated Summary ---

	Department	Avg_Salary	Total_Employees
0	Engineering	88500.0	2
1	HR	61250.0	2
2	Marketing	60000.0	1



2. Creating a Sample DataFrame

PYTHON

Code

```
import pandas as pd

df = pd.DataFrame()

df['Name'] = ['John', 'Emma', 'Liam', 'Olivia']
df['Age'] = [20, 19, 21, 18]
df['Student'] = [True, True, False, True]

print(df)
```

Output

	Name	Age	Student
0	John	20	True
1	Emma	19	True
2	Liam	21	False
3	Olivia	18	True

3. Adding a Row to a DataFrame

PYTHON

Code

```
import pandas as pd

df = pd.DataFrame()
df['Name'] = ['John', 'Emma', 'Liam', 'Olivia']
df['Age'] = [20, 19, 21, 18]
df['Student'] = [True, True, False, True]

new_row = pd.DataFrame([['Sophia', 22, False]], columns=['Name', 'Age', 'Student'])
df = pd.concat([df, new_row], ignore_index=True)
print(df)
```

Output

	Name	Age	Student
0	John	20	True
1	Emma	19	True
2	Liam	21	False
3	Olivia	18	True
4	Sophia	22	False

0	John	20	True
1	Emma	19	True
2	Liam	21	False
3	Olivia	18	True
4	Sophia	22	False

4. Measures of Central Tendency Using SciPy

PYTHON

Code

```
from scipy import stats
import numpy as np

data = [10, 20, 30, 40, 50]
print("Mean:", np.mean(data))
print("Median:", np.median(data))
print("Mode:", stats.mode(data))
```

Output

```
Mean: 30.0
Median: 30.0
Mode: ModeResult(mode=np.int64(10), count=np.int64(1))
```

5. Probability Distribution Analysis Using SciPy

PYTHON

Code

```
from scipy.stats import norm

probability = norm.cdf(85, loc=70, scale=10)
print("Probability:", probability)
```

Output

```
Probability: 0.9331927987311419
```

6. Hypothesis Testing (One-Sample T-Test)

PYTHON

Code

```
from scipy import stats

data = [22, 25, 19, 24, 28, 30]
t_stat, p_value = stats.ttest_1samp(data, 25)
print("T-Statistic:", t_stat)
print("P-Value:", p_value)
```

Output

T-Statistic: -0.2049800154226962
P-Value: 0.8456711339658733

7. Correlation Analysis (Pearson)

PYTHON

Code

```
from scipy.stats import pearsonr

x = [1, 2, 3, 4, 5]
y = [2, 4, 6, 8, 10]
corr, p_value = pearsonr(x, y)
print("Correlation:", corr)
```

Output

Correlation: 1.0

8. Linear Algebra Operations (Solving Linear Equations)

PYTHON

Code

```
from scipy import linalg

A = [[3, 2], [1, 2]]
B = [5, 5]
solution = linalg.solve(A, B)
print(solution)
```

Output

[0. 2.5]

9. Optimization Using SciPy

PYTHON

Code

```
from scipy.optimize import minimize

def objective(x):
    return x**2 + 4

result = minimize(objective, x0=5)
print(result.x)
```

Output

[-2.62955131e-08]

10. Logistic Regression on the Iris Dataset (Scikit-learn)

PYTHON

Code

```
import numpy as np
import matplotlib
matplotlib.use("Agg")
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report

iris = datasets.load_iris()
X = iris.data
y = iris.target

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

log_reg = LogisticRegression()
log_reg.fit(X_train, y_train)

y_pred = log_reg.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

Output

Accuracy: 1.0

11. KMeans Clustering on the Iris Dataset (Scikit-learn)

PYTHON

Code

```
from sklearn.datasets import load_iris
from sklearn.cluster import KMeans

iris = load_iris()
kmeans = KMeans(n_clusters=3, n_init=10, random_state=42)

kmeans.fit(iris.data)
cluster_labels = kmeans.labels_

print("Cluster Labels:", cluster_labels)
```

Output

[illegible]